Pattern Recognition
Homework 4 Announcement

Lastest update: **2023.05.03 12:00**

Homework 4

- Deadline: May. 17, Wed. at 23:59
  - Code assignment (50%)
    - Implement Cross-Validation and Grid Search for SVM training using only NumPy.
  - Questions (50%)
    - Write your answer in detail in the report.
- Question: Link
- Sample code: Link
- Dataset: Link

Support Vector Machines

- The Support Vector Classifier aims to identify the optimal hyperplane for separating distinct classes by maximizing the distance between the sample points and the hyperplane.

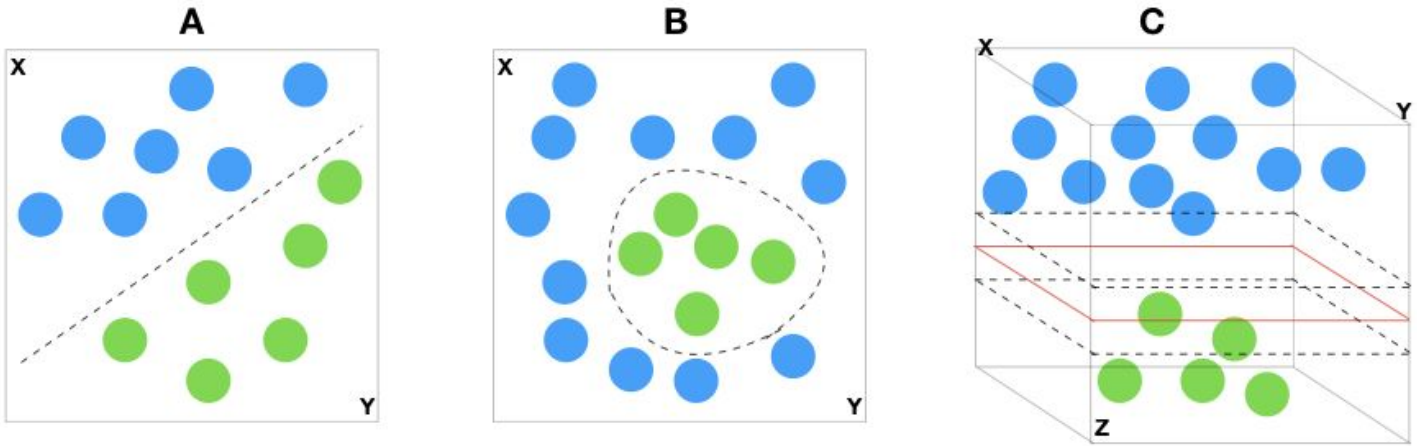

Support Vector Machines

- Since SVM involves complex mathematical operations, you are allowed to use the SVM implementation available in the sklearn library instead of implementing it from scratch.

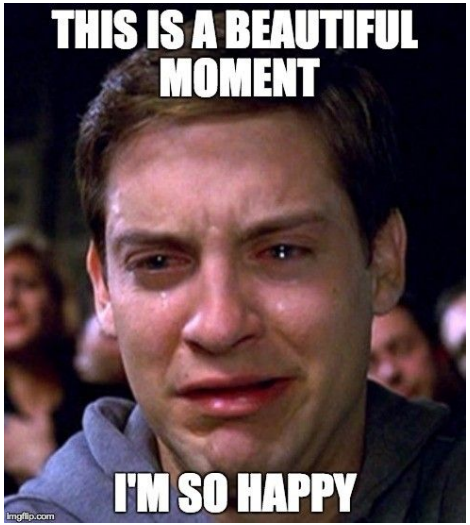

Grid Search and K-Fold Cross-Validation

- There are many hyperparameters in SVM.
- For this homework, you need to perform grid search and cross-validation to find the best hyperparameters for SVM on the provided dataset.

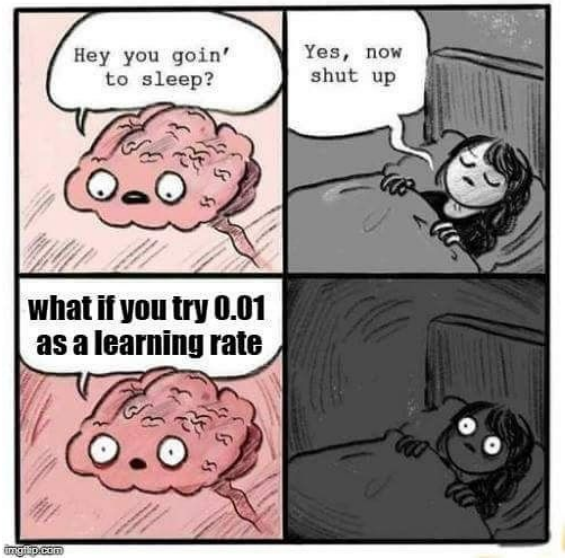

Grid Search

- Suppose we want to find the optimal values of two hyperparameters for an RBF kernel SVM, namely **C** and **gamma**.
- There are numerous combinations that need to be considered!
- Interactive demo

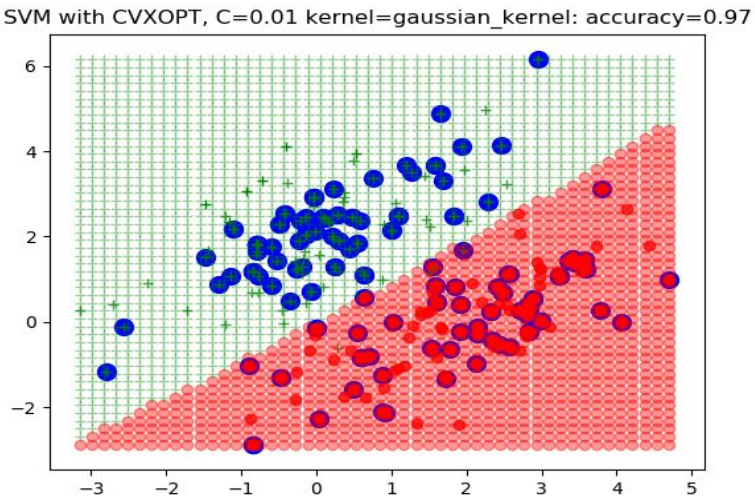

Grid Search

- Grid search explores all possible hyperparameter combinations and selects the best set of hyperparameters based on the model's performance.

```
C = [0.1, 1, 10] #3 values
gamma = [0.01, 0.1, 1, 10] #4 values
# There are totally 12 combinations for tuning
```

C\gamma	0.01	0.1	1	10
0.1	[0.1, 0.01]	[0.1, 0.1]	[0.1, 1]	[0.1, 10]
1	[1, 0.01]	[1, 0.1]	[1, 1]	[1, 10]
10	[10, 0.01]	[10, 0.1]	[10, 1]	[10, 10]

Grid Search

```
Pseudocode
Hyperparameter_One = [a, b, c]
Hyperparameter_Two = [x, y, z]
```

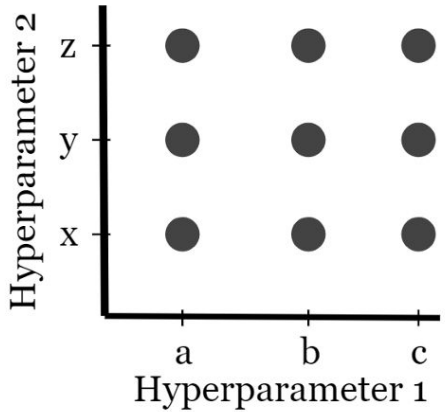

K-Fold Cross-Validation

- We divide the dataset into K subsets.
  - One subset is used for validation.
  - The remaining K-1 subsets are combined to form the training set.
- This process is repeated K times, with each subset used once as the validation data.

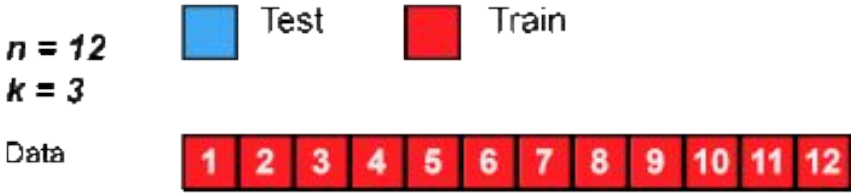

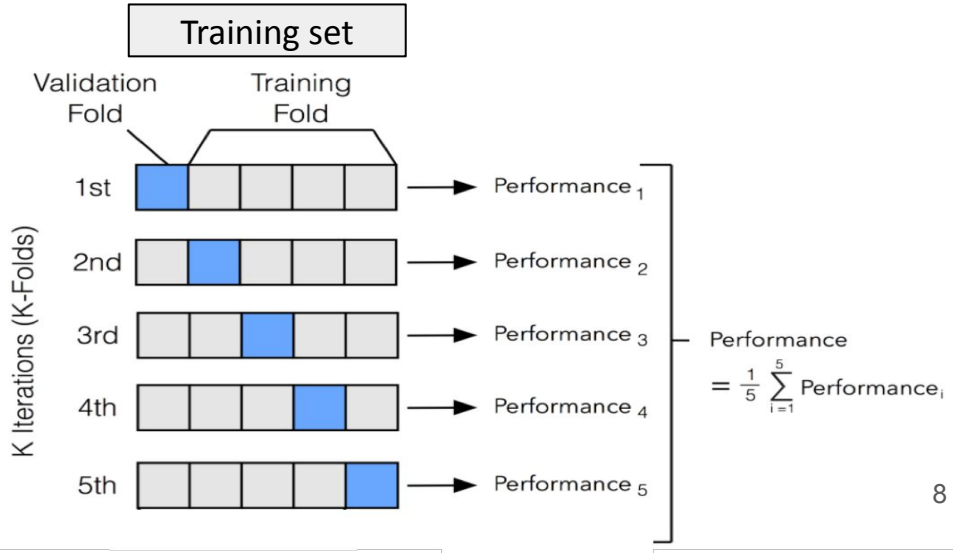

K-Fold Cross-Validation

- Below is the workflow for k-fold cross-validation.

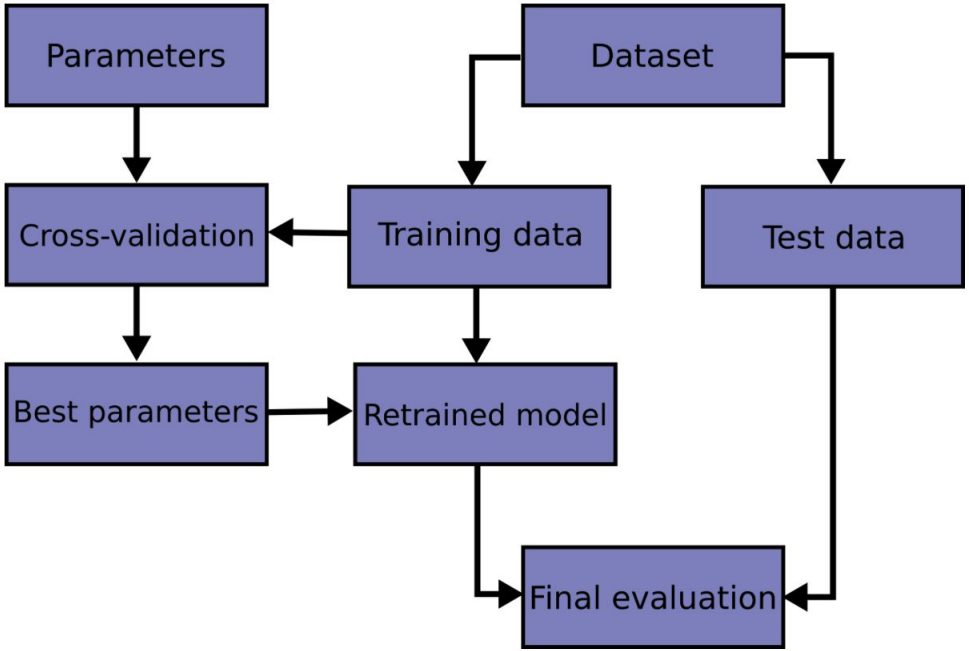

```
C = [0.1, 1, 10] #3 values
gamma = [0.01, 0.1, 1, 10] #4 values
```

Grid Search and K-Fold Cross-Validation

- We can experiment with 12 (3 x 4) combinations of hyperparameters as defined on page 7.
- For each combination, we can apply K-fold cross-validation and calculate the average performance.
- Find the combination that yields the best performance.

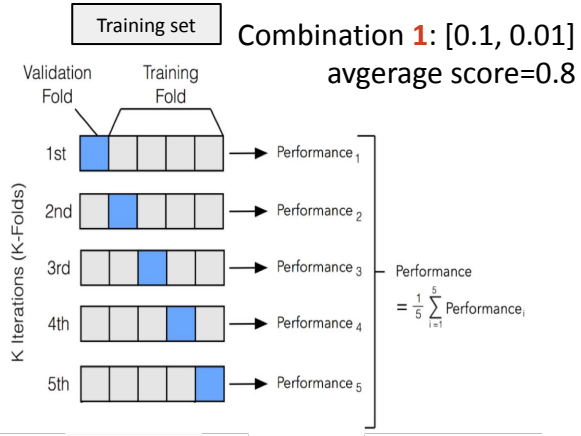

Combination 1: [0.1, 0.01]
avgerage score=0.8

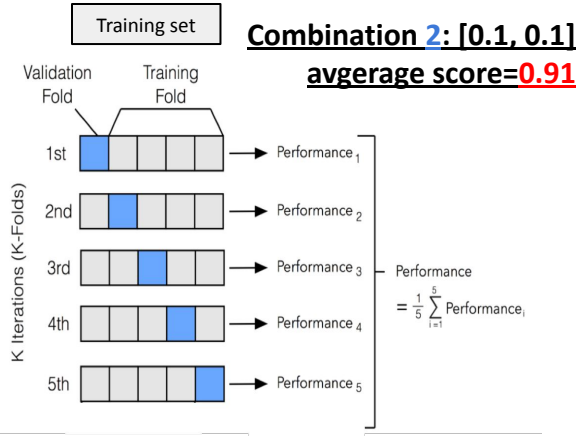

Combination 2: [0.1, 0.1]
avgerage score=0.91

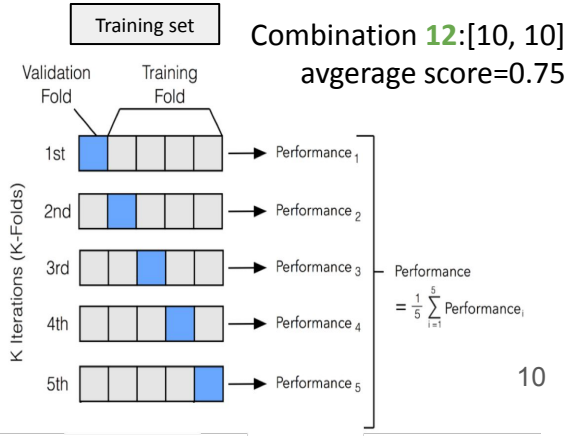

Combination 12:[10, 10]
avgerage score=0.75

Grid Search and K-Fold Cross-Validation

- Finally, train your model on the entire training set using the best hyperparameters, and evaluate the model's performance on the test set.

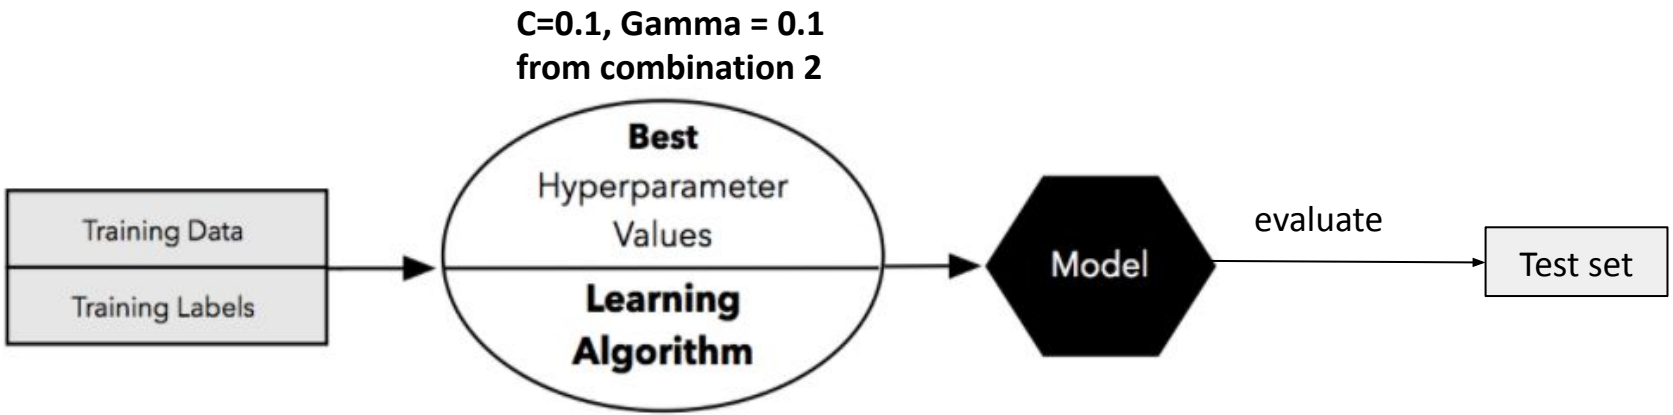

Questions (50%)

1. (10%) Implement K-fold data partitioning using numpy.

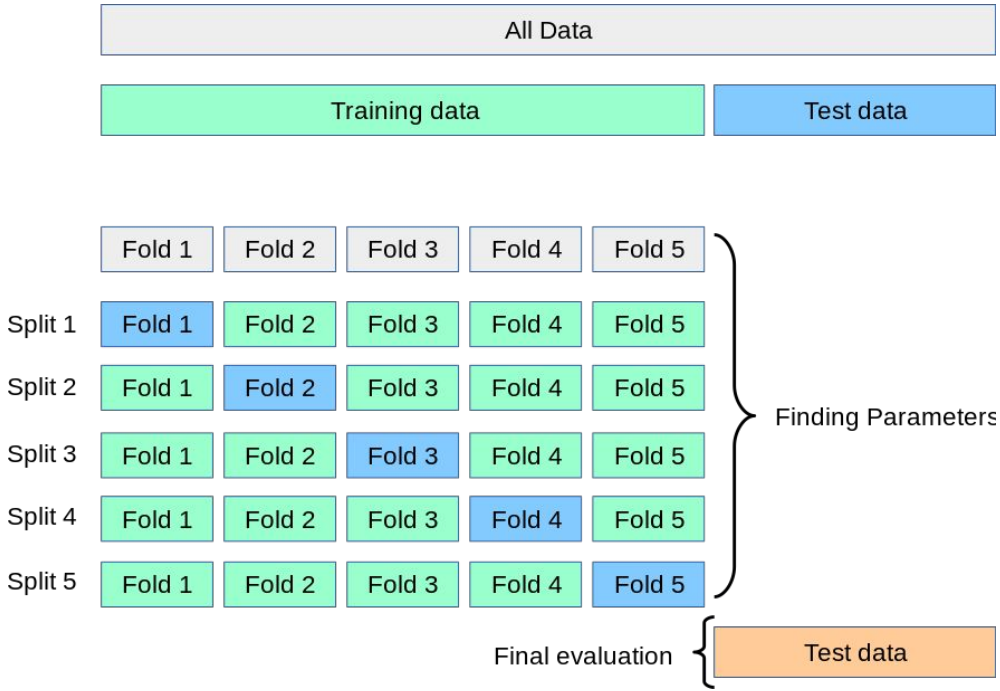

Questions (50%)

2. (10%) Perform a grid search on the hyperparameters **C** and **gamma** to identify the optimal values using cross-validation implemented by NumPy.
3. (10%) Plot the results of your SVM's grid search.

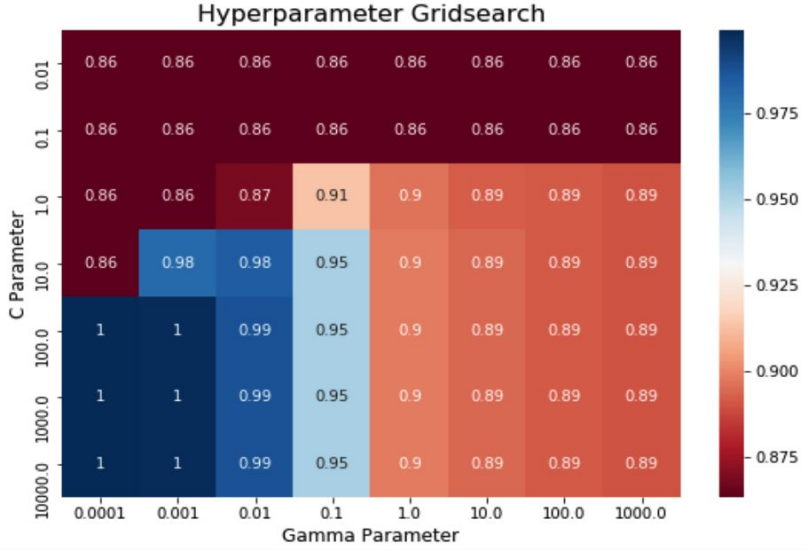

Questions (50%)

4. (20%) Train your SVM model using the best hyperparameters found in Q2 on the entire training dataset, then evaluate its performance on the test set. Print your testing accuracy.

```
# Do Not Modify Below

best_model = SVC(C=best_parameters[0], gamma=best_parameters[1], kernel='rbf')
best_model.fit(x_train, y_train)

y_pred = best_model.predict(x_test)

print("Accuracy score: ", accuracy_score(y_pred, y_test))

# If your accuracy here > 0.9 then you will get full credit (20 points).
```

Accuracy	Your scores
acc > 0.9	20 points
0.85 <= acc <= 0.9	10 points
acc < 0.85	0 points

Dataset

- Training set: 7000
- Testing set: 3000
- 300 features, 2 labels

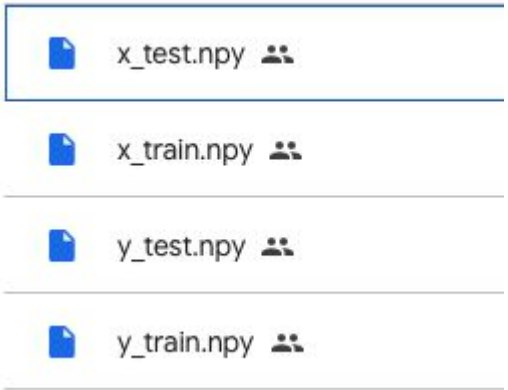

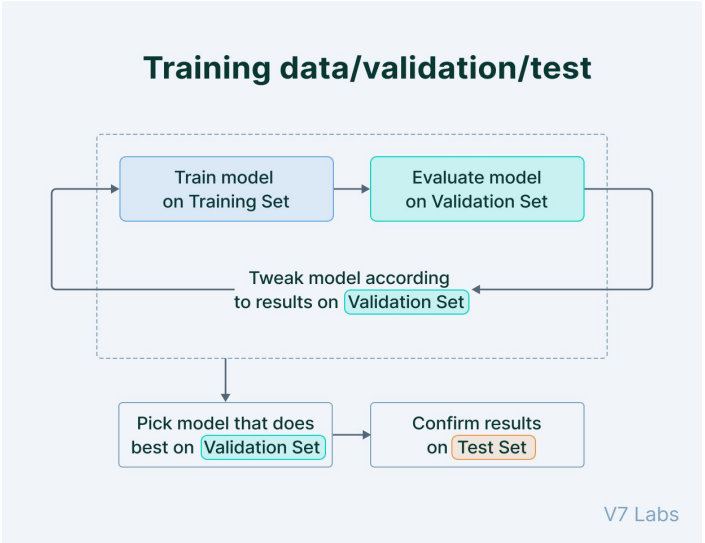

Submission

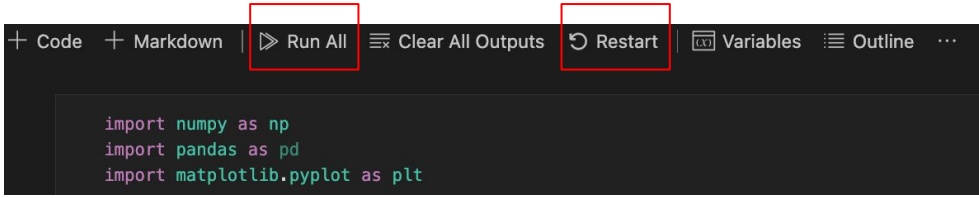

- Compress your .ipynb and .pdf into a zip file and submit it on E3.
- Before submission:
  - Restart and run All
  - Save and submit the .ipynb (keep all cell outputs)
  - Get 0 points if you do not keep the cell outputs.
- <STUDENT ID>_HW4.zip
  - <STUDENT ID>_HW4.ipynb
  - <STUDENT ID>_HW4.pdf
- No csv file in HW4.

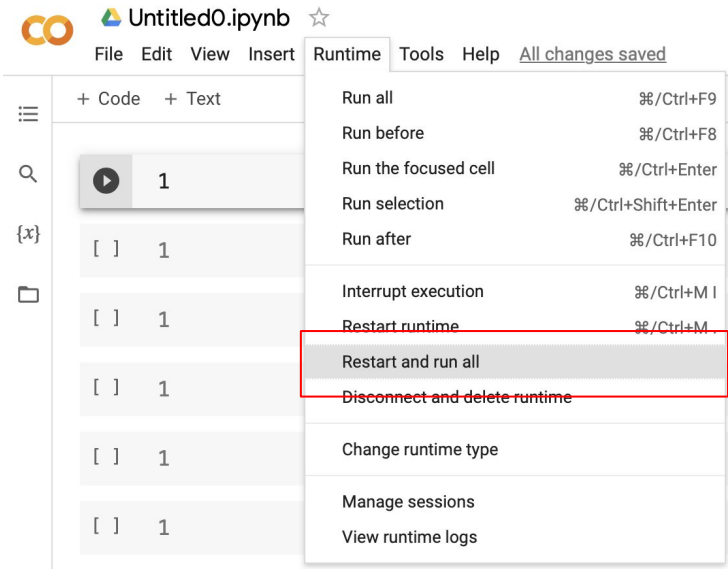


```
❯ zip -r 310551056_HW4.zip 310551056_HW4.ipynb 310551056_HW4.pdf
  adding: 310551056_HW4.ipynb (deflated 37%)
  adding: 310551056_HW4.pdf (deflated 5%)
```

For your reference

Report

- Please write your report in English.
- Please follow the HW1 report template.
- You must type the answer and also screenshot at the same time for the coding part.
- **Answer each question as clearly as possible**. You will get an extra penalty for only the brief answer.

Late policy

- We will deduct a late penalty of 20 points per additional late day.
- If you get 90 points but delay for two days, you get 90 - (20 x 2) = 50 points!
- We only accept submissions that are up to 10 minutes late. Any submissions that are later than that will be considered late, regardless of the reason.

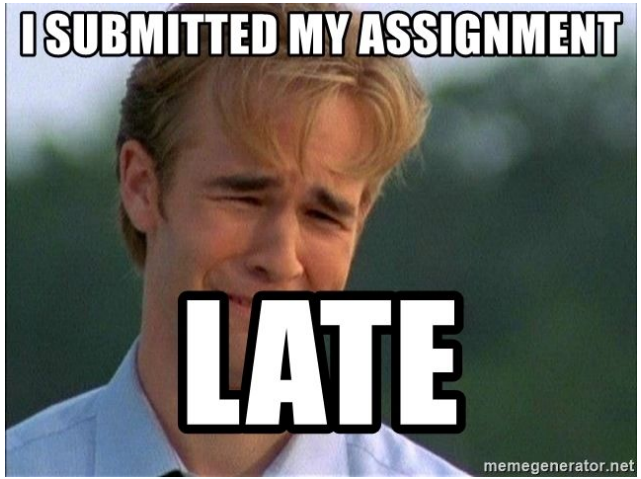

Reference

- K-fold Cross-Validation & Grid Search
- SVM hyperparameter tuning